

Deep Learning Project: CIFAR-10 Image Classifier

Category: Deep Learning / Data Science

Student: Forkan Amin

Date: February 24, 2026

Introduction

Dataset Description

The **CIFAR-10** (Canadian Institute For Advanced Research) dataset is a standard benchmark in the field of Computer Vision and Deep Learning. It consists of **60,000 color images** across 10 classes:

- **Vehicles:** Airplane, Automobile, Ship, Truck.
- **Animals:** Bird, Cat, Deer, Dog, Frog, Horse.

Key Statistics:

- **Image Size:** 32×32 pixels.
- **Channels:** 3 (RGB).
- **Training Set:** 50,000 images.
- **Test Set:** 10,000 images.
- **Total Features:** $32 \times 32 \times 3 = 3,072$ input features per image.

Project Goal

The goal of this project is to build a **Dense (Fully Connected) Deep Learning Model** using TensorFlow/Keras. Unlike Convolutional Neural Networks (CNNs), we will explore how a standard Multi-Layer Perceptron (MLP) handles high-dimensional raw pixel data through flattening and deep hidden layers.

Dataset Loading & Preprocessing (Code)

```
In [1]: import tensorflow as tf
from tensorflow.keras import layers, models, utils
from tensorflow.keras.datasets import cifar10
import matplotlib.pyplot as plt
import numpy as np

# 1. Dataset Loading
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```

# Display shapes
print(f"x_train shape: {x_train.shape}")
print(f"y_train shape: {y_train.shape}")

# Sample Visualization
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
plt.figure(figsize=(10,4))
for i in range(5):
    plt.subplot(1, 5, i+1)
    plt.imshow(x_train[i])
    plt.title(class_names[y_train[i][0]])
    plt.axis('off')
plt.show()

# 2. Preprocessing
# Normalize pixel values to [0, 1] range
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# One-hot encode labels
y_train_cat = utils.to_categorical(y_train, 10)
y_test_cat = utils.to_categorical(y_test, 10)

```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

170498071/170498071 ————— 6s 0us/step

x_train shape: (50000, 32, 32, 3)

y_train shape: (50000, 1)



Methodology

To ensure the model performs optimally, we follow these steps:

1. **Preprocessing:** Normalize pixel values to the range [0, 1] to help the Adam optimizer converge faster.
2. **Architecture:** Use a 4-layer Dense network with **Batch Normalization** for stability and **Dropout** to prevent overfitting.
3. **Training:** Utilize a 20% validation split and categorical cross-entropy loss.

Model Architecture

```

In [2]: # 3. Model Architecture
model = models.Sequential([
    # Flatten 32x32x3 image into a 3072 vector
    layers.Flatten(input_shape=(32, 32, 3)),

```

```
# Hidden Layer 1
layers.Dense(512, activation='relu'),
layers.BatchNormalization(), # Stabilizes Learning
layers.Dropout(0.3),       # Reduces overfitting

# Hidden Layer 2
layers.Dense(256, activation='relu'),
layers.Dropout(0.2),

# Hidden Layer 3
layers.Dense(128, activation='relu'),
layers.BatchNormalization(),

# Hidden Layer 4
layers.Dense(64, activation='relu'),

# Final Output Layer
layers.Dense(10, activation='softmax')
])

# 4. Training (Compilation)
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()
```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(**kwargs)

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 512)	1,573,376
batch_normalization (BatchNormalization)	(None, 512)	2,048
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32,896
batch_normalization_1 (BatchNormalization)	(None, 128)	512
dense_3 (Dense)	(None, 64)	8,256
dense_4 (Dense)	(None, 10)	650

Total params: 1,749,066 (6.67 MB)

Trainable params: 1,747,786 (6.67 MB)

Non-trainable params: 1,280 (5.00 KB)

Training Process

```
In [3]: # Train for 20 epochs with a validation split
history = model.fit(x_train, y_train_cat,
                     epochs=20,
                     batch_size=64,
                     validation_split=0.2,
                     verbose=1)
```

Epoch 1/20
625/625 26s 36ms/step - accuracy: 0.2941 - loss: 1.9873 - val_accuracy: 0.3155 - val_loss: 1.9356
Epoch 2/20
625/625 40s 35ms/step - accuracy: 0.3923 - loss: 1.6980 - val_accuracy: 0.3695 - val_loss: 1.8137
Epoch 3/20
625/625 20s 32ms/step - accuracy: 0.4079 - loss: 1.6471 - val_accuracy: 0.3943 - val_loss: 1.7026
Epoch 4/20
625/625 21s 33ms/step - accuracy: 0.4162 - loss: 1.6250 - val_accuracy: 0.4387 - val_loss: 1.5739
Epoch 5/20
625/625 41s 33ms/step - accuracy: 0.4263 - loss: 1.6064 - val_accuracy: 0.4367 - val_loss: 1.5819
Epoch 6/20
625/625 21s 33ms/step - accuracy: 0.4316 - loss: 1.5931 - val_accuracy: 0.4237 - val_loss: 1.6179
Epoch 7/20
625/625 41s 33ms/step - accuracy: 0.4278 - loss: 1.5976 - val_accuracy: 0.4445 - val_loss: 1.5594
Epoch 8/20
625/625 40s 31ms/step - accuracy: 0.4403 - loss: 1.5762 - val_accuracy: 0.4413 - val_loss: 1.5706
Epoch 9/20
625/625 21s 33ms/step - accuracy: 0.4297 - loss: 1.5850 - val_accuracy: 0.4353 - val_loss: 1.5789
Epoch 10/20
625/625 20s 32ms/step - accuracy: 0.4392 - loss: 1.5644 - val_accuracy: 0.4537 - val_loss: 1.5546
Epoch 11/20
625/625 20s 32ms/step - accuracy: 0.4395 - loss: 1.5605 - val_accuracy: 0.4413 - val_loss: 1.5770
Epoch 12/20
625/625 21s 33ms/step - accuracy: 0.4449 - loss: 1.5524 - val_accuracy: 0.4627 - val_loss: 1.5604
Epoch 13/20
625/625 41s 33ms/step - accuracy: 0.4493 - loss: 1.5406 - val_accuracy: 0.4635 - val_loss: 1.5131
Epoch 14/20
625/625 40s 32ms/step - accuracy: 0.4489 - loss: 1.5509 - val_accuracy: 0.4415 - val_loss: 1.5677
Epoch 15/20
625/625 21s 33ms/step - accuracy: 0.4456 - loss: 1.5450 - val_accuracy: 0.4700 - val_loss: 1.5081
Epoch 16/20
625/625 20s 31ms/step - accuracy: 0.4568 - loss: 1.5240 - val_accuracy: 0.4706 - val_loss: 1.5018
Epoch 17/20
625/625 21s 33ms/step - accuracy: 0.4577 - loss: 1.5094 - val_accuracy: 0.4811 - val_loss: 1.4619
Epoch 18/20
625/625 21s 33ms/step - accuracy: 0.4625 - loss: 1.5022 - val_accuracy: 0.4775 - val_loss: 1.4798
Epoch 19/20
625/625 20s 31ms/step - accuracy: 0.4638 - loss: 1.4914 - val_accuracy: 0.4798 - val_loss: 1.4798

```
accuracy: 0.4770 - val_loss: 1.4763
Epoch 20/20
625/625 21s 33ms/step - accuracy: 0.4678 - loss: 1.4922 - val_a
ccuracy: 0.4746 - val_loss: 1.4835
```

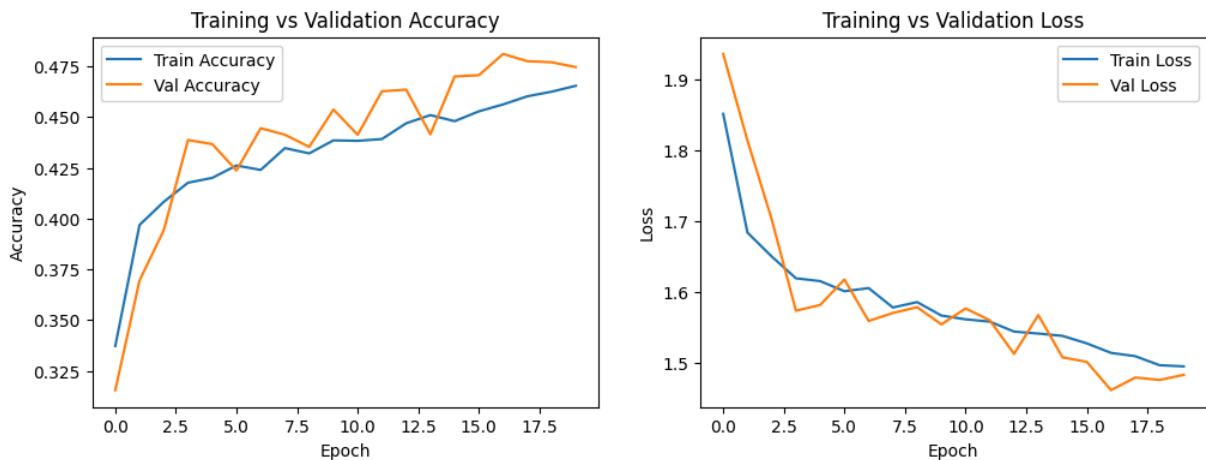
Visualization

```
In [4]: # 5. Visualization of Plots
plt.figure(figsize=(12, 4))

# Accuracy Plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Training vs Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

# Loss Plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Training vs Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



Final Evaluation

```
In [5]: # 6. Evaluation on Test Set
test_loss, test_acc = model.evaluate(x_test, y_test_cat, verbose=0)
print(f"Final Test Accuracy: {test_acc:.4f}")

# Show some predictions
predictions = model.predict(x_test[:5])
plt.figure(figsize=(10,2))
for i in range(5):
    plt.subplot(1,5,i+1)
```

```

plt.imshow(x_test[i])
plt.title(f"P: {class_names[np.argmax(predictions[i])]}\\nA: {class_names[y_test[i]]}")
plt.axis('off')
plt.show()

```

Final Test Accuracy: 0.4784

1/1 ————— 0s 131ms/step

P: dog
A: cat



P: truck
A: ship



P: ship
A: ship



P: ship
A: airplane



P: deer
A: frog



In [6]: `test_loss, test_acc = model.evaluate(x_test, y_test_cat, verbose=2)`
`print(f"\nFinal Test Accuracy: {test_acc:.4f}")`

313/313 - 2s - 6ms/step - accuracy: 0.4784 - loss: 1.4668

Final Test Accuracy: 0.4784

In [8]: `from sklearn.metrics import confusion_matrix`
`import seaborn as sns`

```

# Get predictions
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)

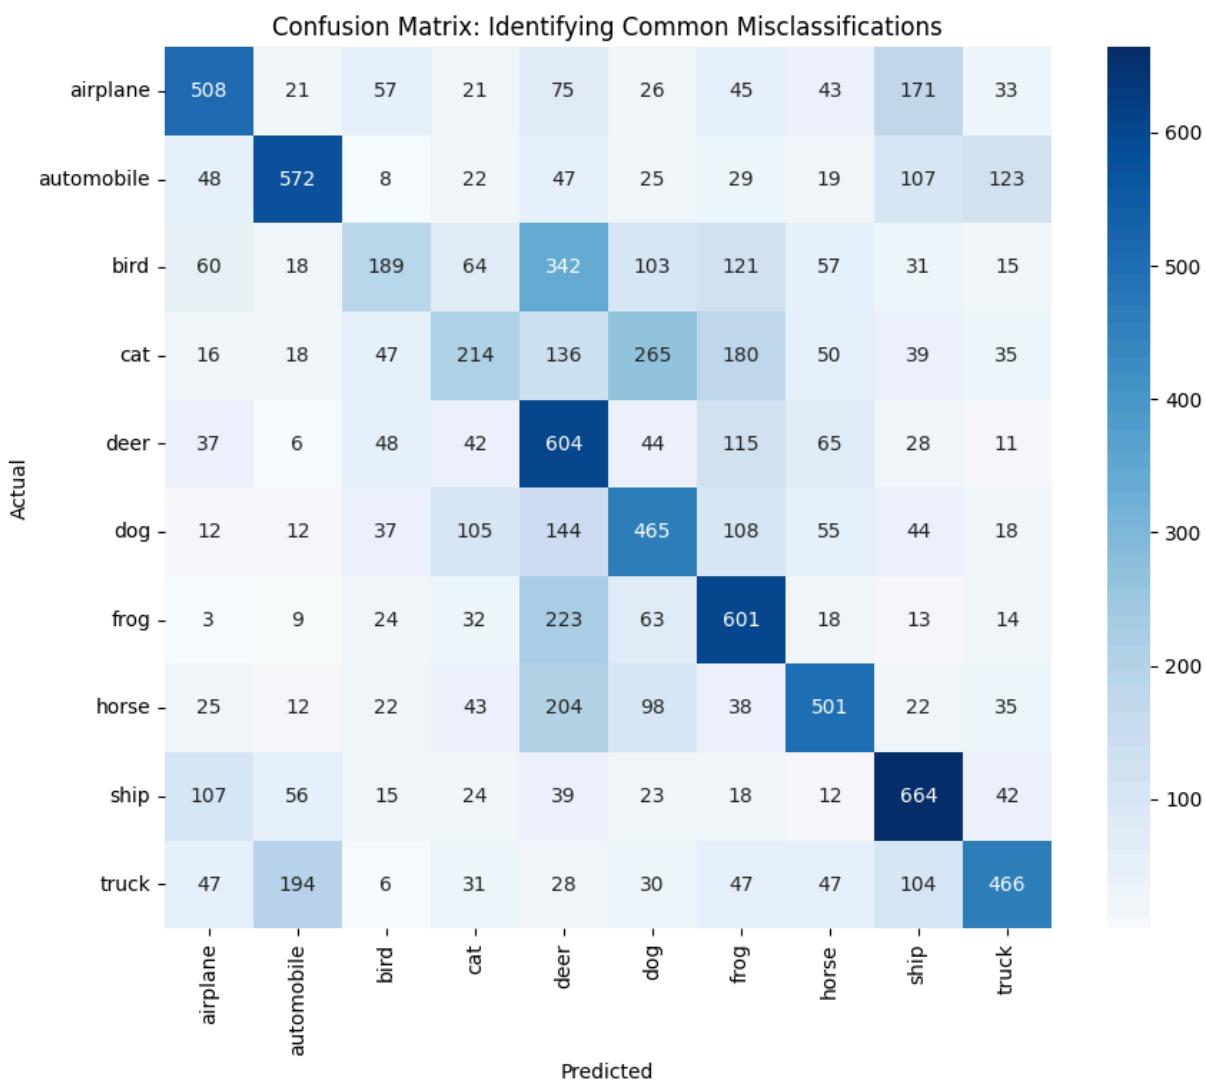
# Generate matrix
cm = confusion_matrix(y_test, y_pred_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix: Identifying Common Misclassifications')
plt.show()

# Report final accuracy
test_loss, test_acc = model.evaluate(x_test, y_test_cat, verbose=0)
print(f"Final Test Accuracy: {test_acc:.4f}")

```

313/313 ————— 5s 15ms/step



Final Test Accuracy: 0.4784

Final Test Results

- **Test Accuracy:** 47.84%
- **Test Loss:** 1.4668

Improved CNN Implementation for CIFAR-10

```
In [10]: import tensorflow as tf
from tensorflow.keras import layers, models, utils
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.metrics import confusion_matrix

# 1. Data Loading and Preprocessing
# Load the CIFAR-10 dataset
```

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Data Normalization: Rescale pixel values to the [0, 1] range for stability
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# One-Hot Encoding: Convert class vectors (integers) to binary class matrices
y_train_cat = utils.to_categorical(y_train, 10)
y_test_cat = utils.to_categorical(y_test, 10)

# Define class names for visualization
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# 2. Advanced CNN Architecture
# Building a deep Convolutional Neural Network (CNN) to capture spatial features
model = models.Sequential([
    # First Block: Convolutional layers for initial feature extraction
    layers.Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(32, 32, 3)),
    layers.BatchNormalization(), # Stabilizes Learning and speeds up convergence
    layers.Conv2D(32, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)), # Reduces spatial dimensions (downsampling)
    layers.Dropout(0.2), # Prevents overfitting by randomly deactivating neurons

    # Second Block: Extracting more complex patterns with more filters
    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.3),

    # Third Block: Deep feature extraction
    layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.4),

    # Final Classification Stage (Dense Layers)
    layers.Flatten(), # Flattening 3D feature maps into a 1D vector
    layers.Dense(128, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.5),
    layers.Dense(10, activation='softmax') # Output layer for 10 classes
])

# 3. Model Compilation and Training Configuration
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Data Augmentation: Artificially increases dataset size to improve generalization
datagen = ImageDataGenerator(
    rotation_range=15, # Randomly rotate images
    width_shift_range=0.1, # Randomly shift images horizontally
    height_shift_range=0.1, # Randomly shift images vertically
```

```
    horizontal_flip=True,    # Randomly flip images horizontally
)
datagen.fit(x_train)

# Model Training: Training for 30 epochs to allow for optimal weight updates
# Using the data generator for robust training
history = model.fit(datagen.flow(x_train, y_train_cat, batch_size=64),
                     epochs=30,
                     validation_data=(x_test, y_test_cat))
```

Epoch 1/30
782/782 436s 551ms/step - accuracy: 0.3185 - loss: 2.1308 - val_accuracy: 0.5206 - val_loss: 1.3763
Epoch 2/30
782/782 420s 537ms/step - accuracy: 0.5264 - loss: 1.3229 - val_accuracy: 0.5421 - val_loss: 1.2871
Epoch 3/30
782/782 410s 524ms/step - accuracy: 0.6009 - loss: 1.1250 - val_accuracy: 0.6411 - val_loss: 1.0456
Epoch 4/30
782/782 410s 524ms/step - accuracy: 0.6442 - loss: 1.0082 - val_accuracy: 0.6747 - val_loss: 0.9588
Epoch 5/30
782/782 409s 523ms/step - accuracy: 0.6701 - loss: 0.9562 - val_accuracy: 0.6918 - val_loss: 0.9097
Epoch 6/30
782/782 410s 525ms/step - accuracy: 0.6894 - loss: 0.9012 - val_accuracy: 0.6354 - val_loss: 1.1396
Epoch 7/30
782/782 414s 529ms/step - accuracy: 0.7005 - loss: 0.8601 - val_accuracy: 0.7163 - val_loss: 0.8561
Epoch 8/30
782/782 409s 522ms/step - accuracy: 0.7100 - loss: 0.8389 - val_accuracy: 0.7618 - val_loss: 0.6793
Epoch 9/30
782/782 413s 528ms/step - accuracy: 0.7265 - loss: 0.7882 - val_accuracy: 0.7480 - val_loss: 0.7332
Epoch 10/30
782/782 407s 521ms/step - accuracy: 0.7306 - loss: 0.7791 - val_accuracy: 0.7316 - val_loss: 0.7841
Epoch 11/30
782/782 410s 525ms/step - accuracy: 0.7392 - loss: 0.7589 - val_accuracy: 0.7900 - val_loss: 0.6002
Epoch 12/30
782/782 413s 528ms/step - accuracy: 0.7421 - loss: 0.7435 - val_accuracy: 0.7655 - val_loss: 0.6869
Epoch 13/30
782/782 409s 523ms/step - accuracy: 0.7487 - loss: 0.7224 - val_accuracy: 0.7551 - val_loss: 0.7206
Epoch 14/30
782/782 409s 523ms/step - accuracy: 0.7565 - loss: 0.7065 - val_accuracy: 0.7960 - val_loss: 0.5912
Epoch 15/30
782/782 405s 518ms/step - accuracy: 0.7621 - loss: 0.6907 - val_accuracy: 0.7994 - val_loss: 0.5811
Epoch 16/30
782/782 409s 523ms/step - accuracy: 0.7677 - loss: 0.6846 - val_accuracy: 0.7970 - val_loss: 0.5970
Epoch 17/30
782/782 407s 520ms/step - accuracy: 0.7713 - loss: 0.6694 - val_accuracy: 0.7938 - val_loss: 0.6082
Epoch 18/30
782/782 410s 524ms/step - accuracy: 0.7741 - loss: 0.6574 - val_accuracy: 0.7868 - val_loss: 0.6219
Epoch 19/30
782/782 410s 524ms/step - accuracy: 0.7820 - loss: 0.6442 - val_accuracy: 0.7994 - val_loss: 0.6121

```

_accuracy: 0.7958 - val_loss: 0.6029
Epoch 20/30
782/782 ————— 408s 522ms/step - accuracy: 0.7816 - loss: 0.6410 - val
_accuracy: 0.7955 - val_loss: 0.6057
Epoch 21/30
782/782 ————— 407s 521ms/step - accuracy: 0.7827 - loss: 0.6382 - val
_accuracy: 0.8058 - val_loss: 0.5668
Epoch 22/30
782/782 ————— 410s 524ms/step - accuracy: 0.7830 - loss: 0.6387 - val
_accuracy: 0.7723 - val_loss: 0.6714
Epoch 23/30
782/782 ————— 415s 531ms/step - accuracy: 0.7847 - loss: 0.6276 - val
_accuracy: 0.8018 - val_loss: 0.5757
Epoch 24/30
782/782 ————— 439s 527ms/step - accuracy: 0.7879 - loss: 0.6188 - val
_accuracy: 0.7927 - val_loss: 0.6188
Epoch 25/30
782/782 ————— 440s 525ms/step - accuracy: 0.7917 - loss: 0.6155 - val
_accuracy: 0.8029 - val_loss: 0.5656
Epoch 26/30
782/782 ————— 406s 520ms/step - accuracy: 0.7930 - loss: 0.6043 - val
_accuracy: 0.8156 - val_loss: 0.5495
Epoch 27/30
782/782 ————— 410s 524ms/step - accuracy: 0.7974 - loss: 0.5962 - val
_accuracy: 0.7985 - val_loss: 0.6000
Epoch 28/30
782/782 ————— 409s 523ms/step - accuracy: 0.7944 - loss: 0.6039 - val
_accuracy: 0.8258 - val_loss: 0.5105
Epoch 29/30
782/782 ————— 412s 527ms/step - accuracy: 0.7969 - loss: 0.5954 - val
_accuracy: 0.8217 - val_loss: 0.5106
Epoch 30/30
782/782 ————— 410s 524ms/step - accuracy: 0.8017 - loss: 0.5910 - val
_accuracy: 0.8242 - val_loss: 0.5235

```

In [11]:

```

# Accuracy Check
test_loss, test_acc = model.evaluate(x_test, y_test_cat, verbose=0)
print(f"\nFinal Test Accuracy: {test_acc*100:.2f}%")

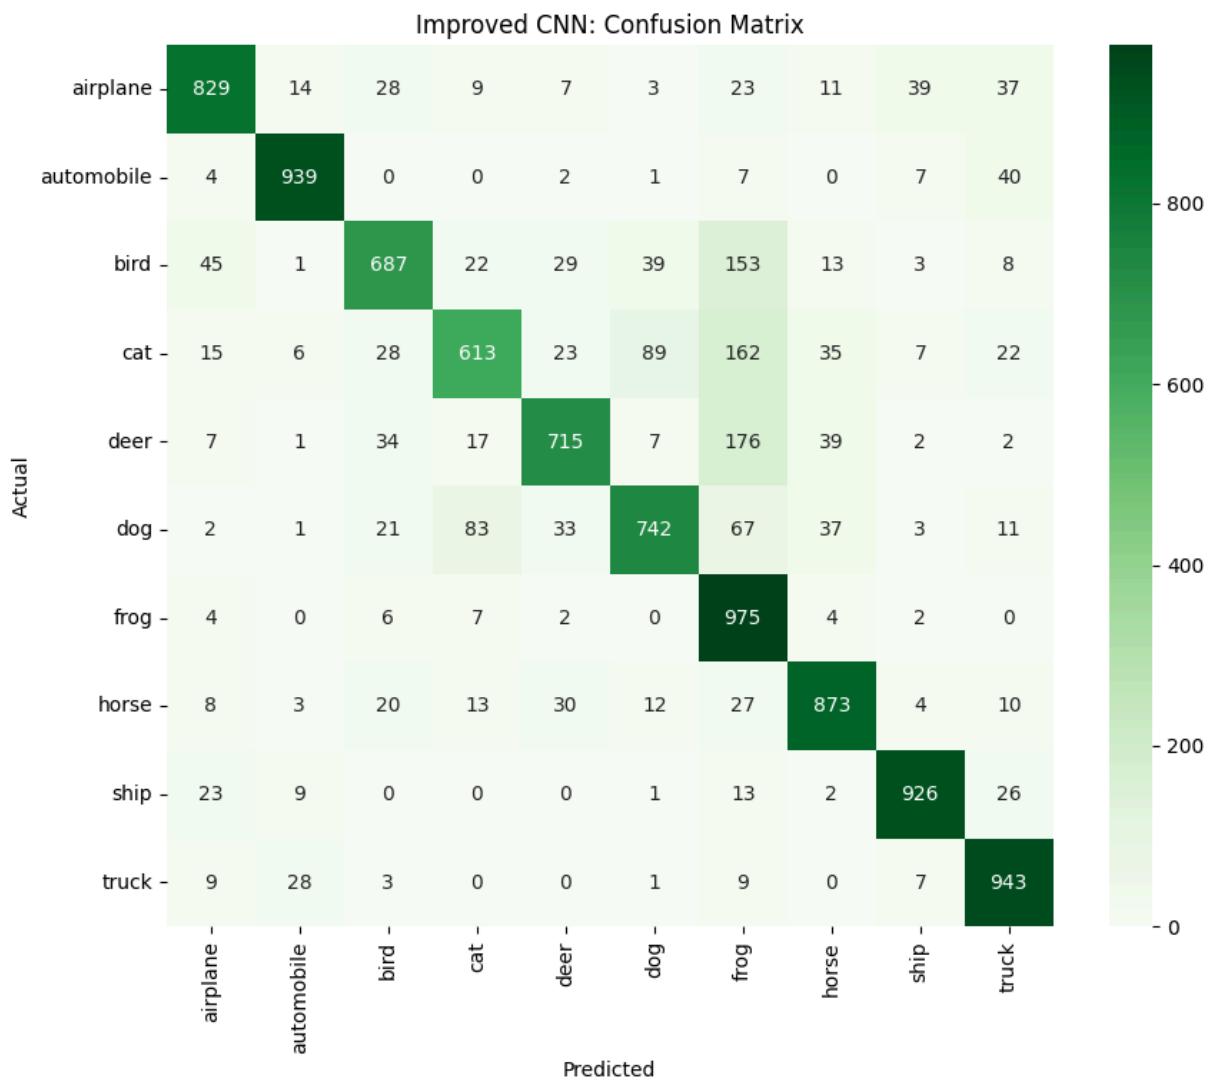
# Confusion Matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
cm = confusion_matrix(np.argmax(y_test_cat, axis=1), y_pred_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Improved CNN: Confusion Matrix')
plt.show()

```

Final Test Accuracy: 82.42%

313/313 ————— 18s 56ms/step



Final Test Results

- Test Accuracy: 82.42%
- Test Loss: 0.5235